

DESeq Mini Project: Pathway Analysis from RNA-Seq Results

Andrew Kapinos

11/19/2021

Section 1. Differential Expression Analysis

```
library(DESeq2)
```

```
## Warning: package 'GenomicRanges' was built under R version 4.1.2
```

Let's load our count and metagene data files. We can then import our metadata.

```
metaFile <- "GSE37704_metadata.csv"
countFile <- "GSE37704_featurecounts.csv"

# Import metadata and take a look
colData = read.csv(metaFile, row.names=1)
head(colData)
```

```
##              condition
## SRR493366 control_sirna
## SRR493367 control_sirna
## SRR493368 control_sirna
## SRR493369      hoxa1_kd
## SRR493370      hoxa1_kd
## SRR493371      hoxa1_kd
```

Let's import our count data as well.

```
countData.raw = read.csv(countFile, row.names=1)
head(countData.raw)
```

```
##              length SRR493366 SRR493367 SRR493368 SRR493369 SRR493370
## ENSG00000186092    918         0         0         0         0         0
## ENSG00000279928    718         0         0         0         0         0
## ENSG00000279457   1982        23        28        29        29        28
## ENSG00000278566    939         0         0         0         0         0
## ENSG00000273547    939         0         0         0         0         0
## ENSG00000187634   3214        124        123        205        207        212
##              SRR493371
## ENSG00000186092         0
## ENSG00000279928         0
```

```
## ENSG00000279457      46
## ENSG00000278566       0
## ENSG00000273547       0
## ENSG00000187634     258
```

Q. Complete the code below to remove the troublesome first column from countData.

```
countData <- as.matrix(countData.raw[, -1])
head(countData)
```

```
##           SRR493366 SRR493367 SRR493368 SRR493369 SRR493370 SRR493371
## ENSG00000186092      0         0         0         0         0         0
## ENSG00000279928      0         0         0         0         0         0
## ENSG00000279457     23        28        29        29        28        46
## ENSG00000278566      0         0         0         0         0         0
## ENSG00000273547      0         0         0         0         0         0
## ENSG00000187634    124       123       205       207       212       258
```

Q. Complete the code below to filter countData to exclude genes (i.e. rows) where we have 0 read count across all samples (i.e. columns).

Tip: What will rowSums() of countData return and how could you use it in this context?

rowSums(countData)==0 will return the row numbers of rows with all 0s, which we want to remove. We can add the minus symbol to select rows that don't have all 0s.

```
# Filter count data where you have 0 read count across all samples.
countData = countData[-which(rowSums(countData)==0),]
head(countData)
```

```
##           SRR493366 SRR493367 SRR493368 SRR493369 SRR493370 SRR493371
## ENSG00000279457     23        28        29        29        28        46
## ENSG00000187634    124       123       205       207       212       258
## ENSG00000188976   1637      1831      2383      1226      1326      1504
## ENSG00000187961    120       153       180       236       255       357
## ENSG00000187583     24        48        65        44        48        64
## ENSG00000187642      4         9        16        14        16        16
```

Optional: PCA Analysis

Let's perform PCA to ensure that the data looks well separated.

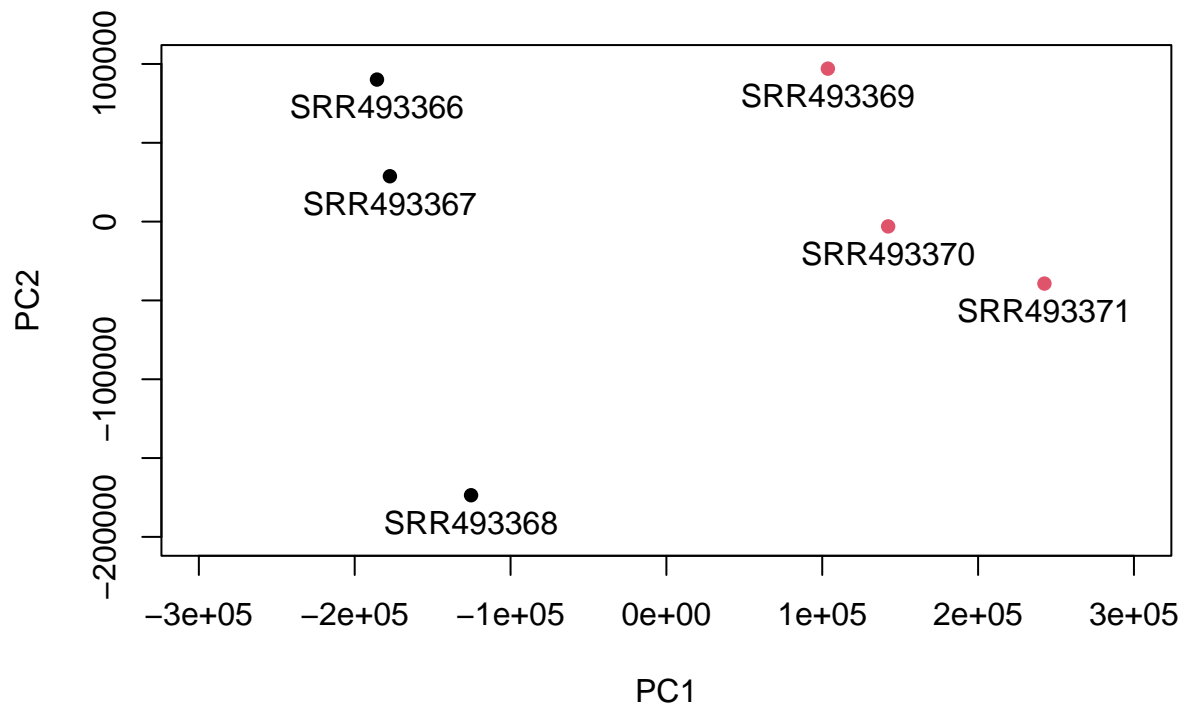
```
pca.counts <- prcomp(t(countData))
attributes(pca.counts)
```

```
## $names
## [1] "sdev"      "rotation" "center"    "scale"     "x"
##
## $class
## [1] "prcomp"
```

```
summary(pca.counts$x)
```

```
##      PC1      PC2      PC3      PC4
## Min.   :-185705 Min.   :-173585 Min.   :-33499 Min.   :-10518
## 1st Qu.: -164392 1st Qu.: -30247  1st Qu.: -3696  1st Qu.: -4387
## Median : -10861  Median :  12872  Median : -1112  Median :  2204
## Mean    :      0   Mean    :      0   Mean    :      0   Mean    :      0
## 3rd Qu.:  132618 3rd Qu.:  74778  3rd Qu.: 10853  3rd Qu.:  5270
## Max.    :  242550 Max.    :  97047  Max.    : 25439  Max.    :  6403
##      PC5      PC6
## Min.   :-6734.7 Min.   :-3.045e-09
## 1st Qu.: -3527.9 1st Qu.: -4.769e-10
## Median :   323.4 Median :  1.729e-10
## Mean    :    0.0 Mean    : -1.890e-12
## 3rd Qu.:  2729.8 3rd Qu.:  1.178e-09
## Max.    :  7367.7 Max.    :  1.868e-09
```

```
plot(pca.counts$x[,1], pca.counts$x[,2],
     xlim=c(-300000,300000), ylim=c(-200000,100000),
     xlab="PC1",
     ylab="PC2", col=as.factor(colData$condition), pch=16)
text(pca.counts$x[,1], pca.counts$x[,2], colnames(countData), pos=1)
```



It looks like the data has some trends between the two groups based on their separation along PC1 in the PC plot. Let's continue to DESeq2.

Running DESeq2

Q. Call the `summary()` function on your results to get a sense of how many genes are up or down-regulated at the default 0.1 p-value cutoff.

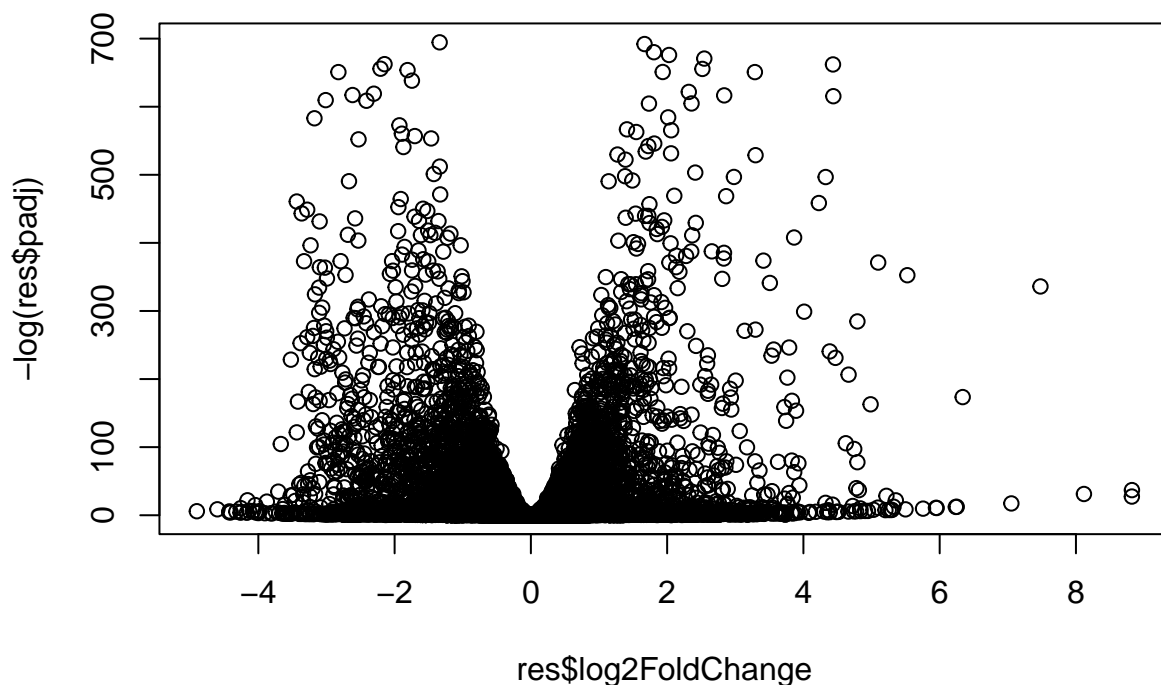
```
dds = DESeqDataSetFromMatrix(countData=countData,  
                              colData=colData,  
                              design=~condition)  
  
dds = DESeq(dds)  
res = results(dds)  
summary(res)
```

```
##  
## out of 15975 with nonzero total read count  
## adjusted p-value < 0.1  
## LFC > 0 (up)      : 4349, 27%  
## LFC < 0 (down)    : 4396, 28%  
## outliers [1]      : 0, 0%  
## low counts [2]    : 1237, 7.7%  
## (mean count < 0)  
## [1] see 'cooksCutoff' argument of ?results  
## [2] see 'independentFiltering' argument of ?results
```

27% are upregulated and 28% are downregulated.

Volcano plot

```
plot(res$log2FoldChange, -log(res$padj))
```



Q. Improve this plot by completing the below code, which adds color and axis labels

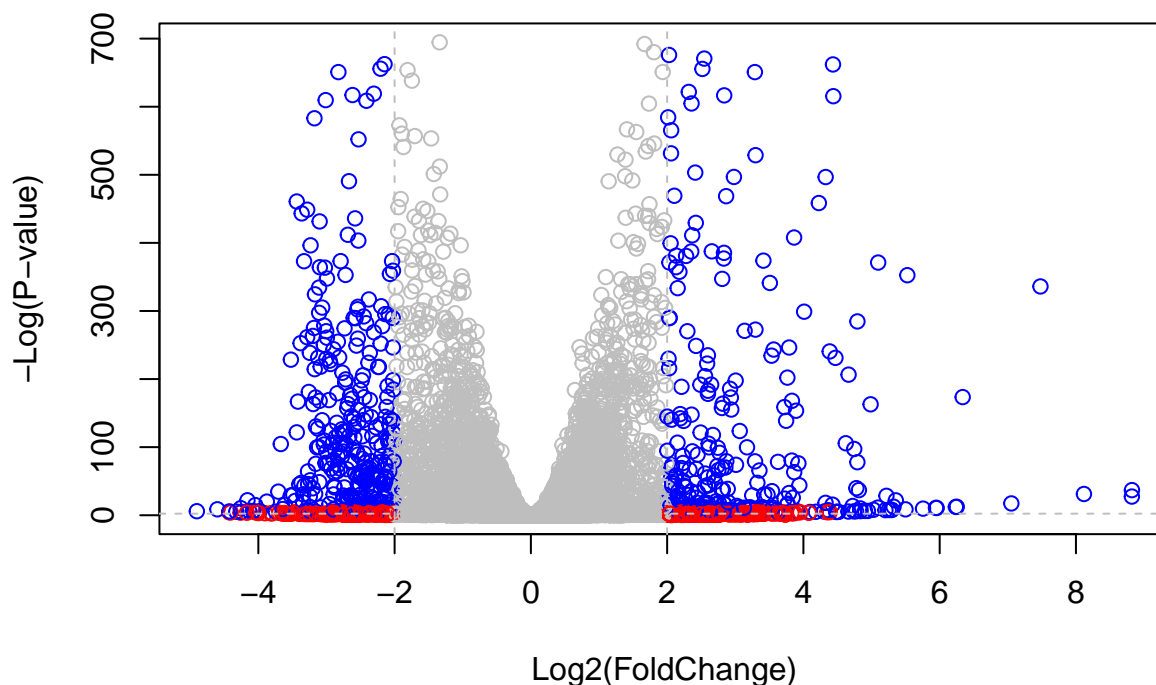
```
# Make a color vector for all genes
mycols <- rep("gray", nrow(res) )

# Color red the genes with absolute fold change above 2
mycols[ abs(res$log2FoldChange) > 2 ] <- "red"

# Color blue those with adjusted p-value less than 0.01
# and absolute fold change more than 2
inds <- (res$padj < 0.01) & (abs(res$log2FoldChange) > 2 )
mycols[ inds ] <- "blue"

plot( res$log2FoldChange, -log(res$padj), col=mycols, xlab="Log2(FoldChange)", ylab="-Log(P-value)" )

# Cut-off lines
abline(v=c(-2,2), col="gray", lty=2)
abline(h=-log(0.1), col="gray", lty=2)
```



Adding gene annotation

Q. Use the `mapIds()` function multiple times to add SYMBOL, ENTREZID and GENENAME annotation to our results by completing the code below.

```
library("AnnotationDbi")
```

```
## Warning: package 'AnnotationDbi' was built under R version 4.1.2
```

```
library("org.Hs.eg.db")
```

```
columns(org.Hs.eg.db)
```

```
## [1] "ACCNUM"      "ALIAS"       "ENSEMBL"     "ENSEMBLPROT" "ENSEMBLTRANS"
## [6] "ENTREZID"    "ENZYME"      "EVIDENCE"     "EVIDENCEALL"  "GENENAME"
## [11] "GENETYPE"    "GO"          "GOALL"        "IPI"          "MAP"
## [16] "OMIM"        "ONTOLOGY"    "ONTOLOGYALL"  "PATH"         "PFAM"
## [21] "PMID"        "PROSITE"     "REFSEQ"       "SYMBOL"       "UCSCCKG"
## [26] "UNIPROT"
```

```
res$symbol = mapIds(org.Hs.eg.db,
                    keys=row.names(res),
                    keytype="ENSEMBL",
```

```

        column="SYMBOL",
        multiVals="first")

res$entrez = mapIds(org.Hs.eg.db,
                    keys=row.names(res),
                    keytype="ENSEMBL",
                    column="ENTREZID",
                    multiVals="first")

res$name = mapIds(org.Hs.eg.db,
                  keys=row.names(res),
                  keytype="ENSEMBL",
                  column="GENENAME",
                  multiVals="first")

head(res, 10)

```

```

## log2 fold change (MLE): condition hoxa1 kd vs control sirna
## Wald test p-value: condition hoxa1 kd vs control sirna
## DataFrame with 10 rows and 9 columns
##
```

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
## ENSG00000279457	29.913579	0.1792571	0.3248216	0.551863	5.81042e-01
## ENSG00000187634	183.229650	0.4264571	0.1402658	3.040350	2.36304e-03
## ENSG00000188976	1651.188076	-0.6927205	0.0548465	-12.630158	1.43990e-36
## ENSG00000187961	209.637938	0.7297556	0.1318599	5.534326	3.12428e-08
## ENSG00000187583	47.255123	0.0405765	0.2718928	0.149237	8.81366e-01
## ENSG00000187642	11.979750	0.5428105	0.5215598	1.040744	2.97994e-01
## ENSG00000188290	108.922128	2.0570638	0.1969053	10.446970	1.51282e-25
## ENSG00000187608	350.716868	0.2573837	0.1027266	2.505522	1.22271e-02
## ENSG00000188157	9128.439422	0.3899088	0.0467163	8.346304	7.04321e-17
## ENSG00000237330	0.158192	0.7859552	4.0804729	0.192614	8.47261e-01

```

##
```

	padj	symbol	entrez	name
	<numeric>	<character>	<character>	<character>
## ENSG00000279457	6.86555e-01	WASH9P	102723897	WAS protein family h..
## ENSG00000187634	5.15718e-03	SAMD11	148398	sterile alpha motif ..
## ENSG00000188976	1.76549e-35	NOC2L	26155	NOC2 like nucleolar ..
## ENSG00000187961	1.13413e-07	KLHL17	339451	kelch like family me..
## ENSG00000187583	9.19031e-01	PLEKHN1	84069	pleckstrin homology ..
## ENSG00000187642	4.03379e-01	PERM1	84808	PPARGC1 and ESRR ind..
## ENSG00000188290	1.30538e-24	HES4	57801	hes family bHLH tran..
## ENSG00000187608	2.37452e-02	ISG15	9636	ISG15 ubiquitin like..
## ENSG00000188157	4.21963e-16	AGRN	375790	agrin
## ENSG00000237330	NA	RNF223	401934	ring finger protein ..

Q. Finally for this section let's reorder these results by adjusted p-value and save them to a CSV file in your current project directory.

```

res = res[order(res$pvalue),]
write.csv(res, file="deseq_results.csv")

```

Section 2: Pathway Analysis

Let's load the packages and data we'll need.

```
library(pathview)
```

```
## #####  
## Pathview is an open source software package distributed under GNU General  
## Public License version 3 (GPLv3). Details of GPLv3 is available at  
## http://www.gnu.org/licenses/gpl-3.0.html. Particullary, users are required to  
## formally cite the original Pathview paper (not just mention it) in publications  
## or products. For details, do citation("pathview") within R.  
##  
## The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG  
## license agreement (details at http://www.kegg.jp/kegg/legal.html).  
## #####
```

```
library(gage)
```

```
##
```

```
library(gageData)
```

```
data(kegg.sets.hs)
```

```
data(sigmet.idx.hs)
```

```
# Focus on signaling and metabolic pathways only
```

```
kegg.sets.hs = kegg.sets.hs[sigmet.idx.hs]
```

```
# Examine the first 3 pathways
```

```
head(kegg.sets.hs, 3)
```

```
## $'hsa00232 Caffeine metabolism'
```

```
## [1] "10" "1544" "1548" "1549" "1553" "7498" "9"
```

```
##
```

```
## $'hsa00983 Drug metabolism - other enzymes'
```

```
## [1] "10" "1066" "10720" "10941" "151531" "1548" "1549" "1551"
```

```
## [9] "1553" "1576" "1577" "1806" "1807" "1890" "221223" "2990"
```

```
## [17] "3251" "3614" "3615" "3704" "51733" "54490" "54575" "54576"
```

```
## [25] "54577" "54578" "54579" "54600" "54657" "54658" "54659" "54963"
```

```
## [33] "574537" "64816" "7083" "7084" "7172" "7363" "7364" "7365"
```

```
## [41] "7366" "7367" "7371" "7372" "7378" "7498" "79799" "83549"
```

```
## [49] "8824" "8833" "9" "978"
```

```
##
```

```
## $'hsa00230 Purine metabolism'
```

```
## [1] "100" "10201" "10606" "10621" "10622" "10623" "107" "10714"
```

```
## [9] "108" "10846" "109" "111" "11128" "11164" "112" "113"
```

```
## [17] "114" "115" "122481" "122622" "124583" "132" "158" "159"
```

```
## [25] "1633" "171568" "1716" "196883" "203" "204" "205" "221823"
```

```
## [33] "2272" "22978" "23649" "246721" "25885" "2618" "26289" "270"
```

```
## [41] "271" "27115" "272" "2766" "2977" "2982" "2983" "2984"
```

```
## [49] "2986" "2987" "29922" "3000" "30833" "30834" "318" "3251"
```



```
## [57] "353"      "3614"     "3615"     "3704"     "377841"   "471"      "4830"     "4831"
## [65] "4832"     "4833"     "4860"     "4881"     "4882"     "4907"     "50484"    "50940"
## [73] "51082"    "51251"    "51292"    "5136"     "5137"     "5138"     "5139"     "5140"
## [81] "5141"     "5142"     "5143"     "5144"     "5145"     "5146"     "5147"     "5148"
## [89] "5149"     "5150"     "5151"     "5152"     "5153"     "5158"     "5167"     "5169"
## [97] "51728"    "5198"     "5236"     "5313"     "5315"     "53343"    "54107"    "5422"
## [105] "5424"     "5425"     "5426"     "5427"     "5430"     "5431"     "5432"     "5433"
## [113] "5434"     "5435"     "5436"     "5437"     "5438"     "5439"     "5440"     "5441"
## [121] "5471"     "548644"   "55276"    "5557"     "5558"     "55703"    "55811"    "55821"
## [129] "5631"     "5634"     "56655"    "56953"    "56985"    "57804"    "58497"    "6240"
## [137] "6241"     "64425"    "646625"   "654364"   "661"      "7498"     "8382"     "84172"
## [145] "84265"    "84284"    "84618"    "8622"     "8654"     "87178"    "8833"     "9060"
## [153] "9061"     "93034"    "953"      "9533"     "954"      "955"      "956"      "957"
## [161] "9583"     "9615"
```

To use `gage()` we'll need a named vector of fold changes.

```
foldchanges = res$log2FoldChange
names(foldchanges) = res$entrez
head(foldchanges)
```

```
##      1266      54855      1465      51232      2034      2317
## -2.422719  3.201955 -2.313738 -2.059631 -1.888019 -1.649792
```

Now let's run the `gage` pathway analysis.

```
keggres = gage(foldchanges, gsets=kegg.sets.hs)
```

Let's take a look at the result structure and the first few down (less) pathways.

```
attributes(keggres)
```

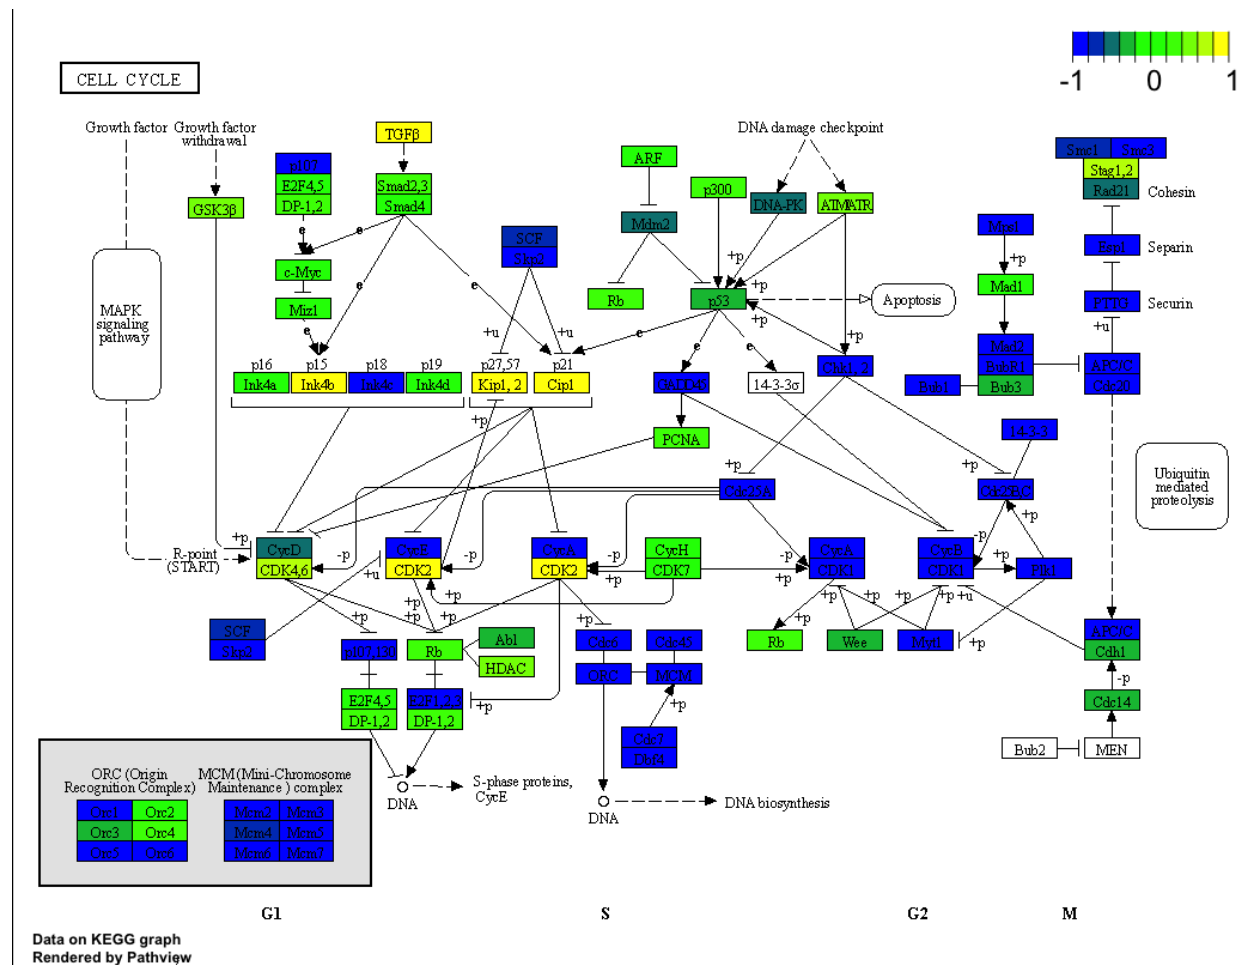
```
## $names
## [1] "greater" "less"    "stats"
```

```
head(keggres$less)
```

```
##                p.geomean stat.mean      p.val
## hsa04110 Cell cycle      8.995727e-06 -4.378644 8.995727e-06
## hsa03030 DNA replication  9.424076e-05 -3.951803 9.424076e-05
## hsa03013 RNA transport    1.375901e-03 -3.028500 1.375901e-03
## hsa03440 Homologous recombination 3.066756e-03 -2.852899 3.066756e-03
## hsa04114 Oocyte meiosis    3.784520e-03 -2.698128 3.784520e-03
## hsa00010 Glycolysis / Gluconeogenesis 8.961413e-03 -2.405398 8.961413e-03
##                q.val set.size      exp1
## hsa04110 Cell cycle    0.001448312      121 8.995727e-06
## hsa03030 DNA replication 0.007586381       36 9.424076e-05
## hsa03013 RNA transport  0.073840037      144 1.375901e-03
## hsa03440 Homologous recombination 0.121861535       28 3.066756e-03
## hsa04114 Oocyte meiosis 0.121861535      102 3.784520e-03
## hsa00010 Glycolysis / Gluconeogenesis 0.212222694       53 8.961413e-03
```

And let's generate a pathway figure from KEGG data using our results.

```
pathview(gene.data=foldchanges, pathway.id="hsa04110", low="blue", mid="green", high="yellow")
```



Let's take a look at the up (more) pathways, too.

```
## Focus on top 5 upregulated pathways here for demo purposes only
```

```
keggrespathways <- rownames(keggres$greater)[1:5]
```

```
# Extract the 8 character long IDs part of each string
```

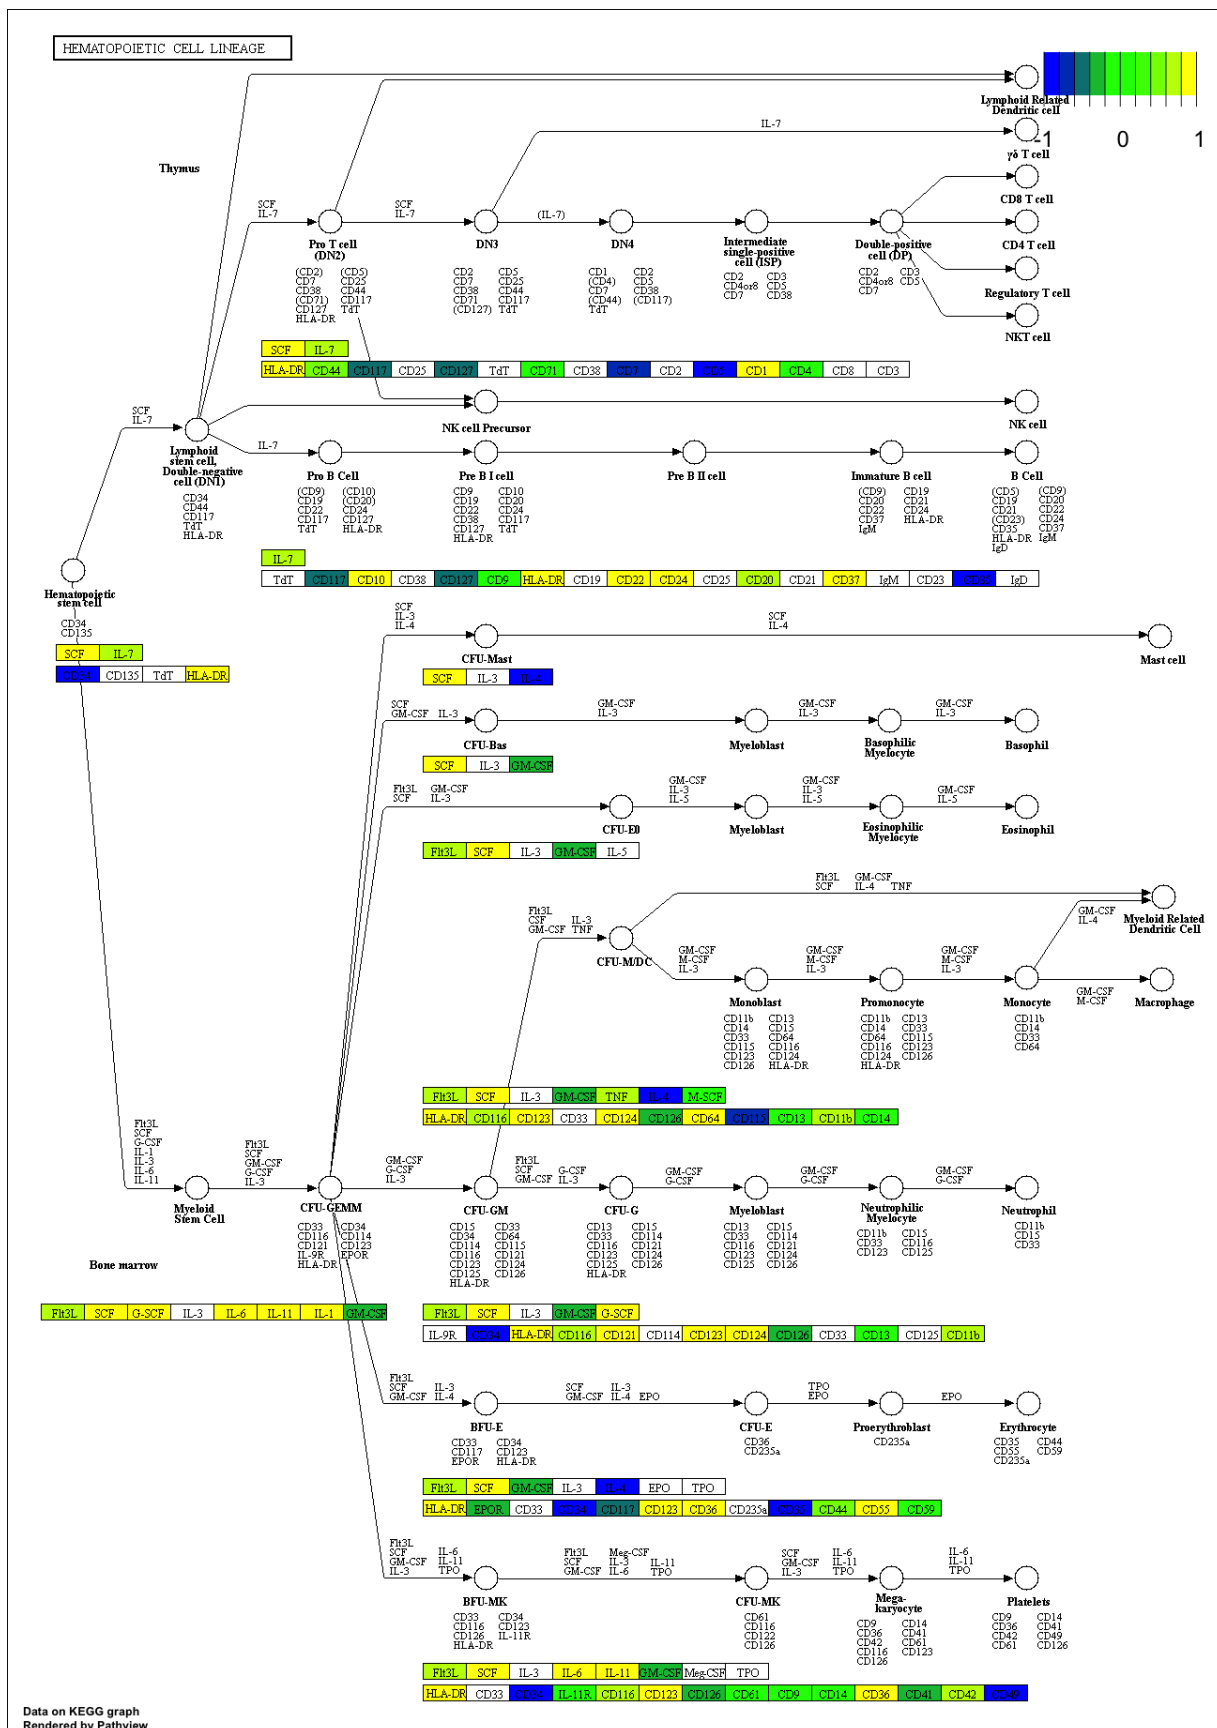
```
keggresids = substr(keggrespathways, start=1, stop=8)
```

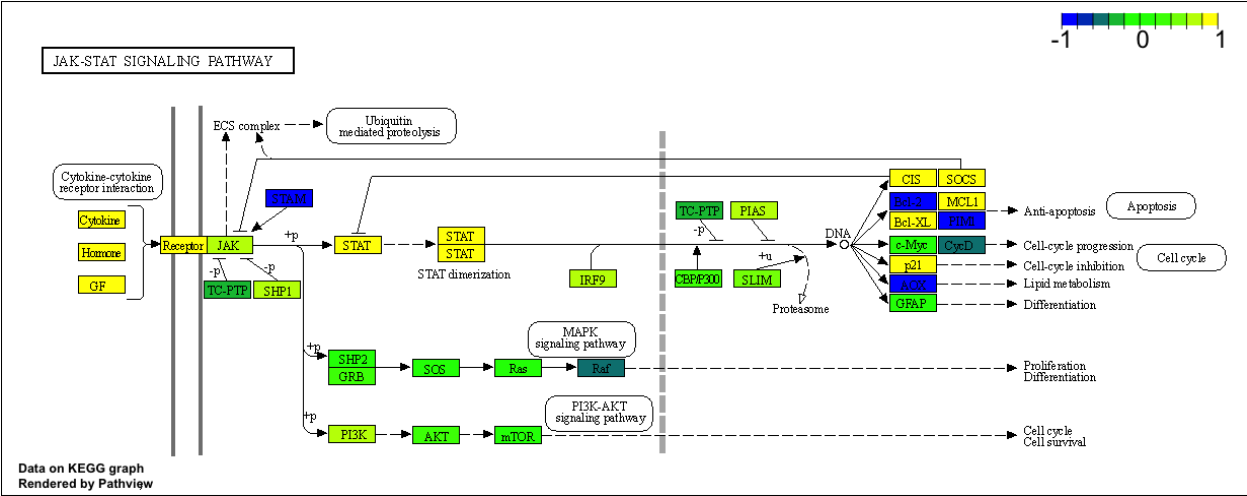
```
keggresids
```

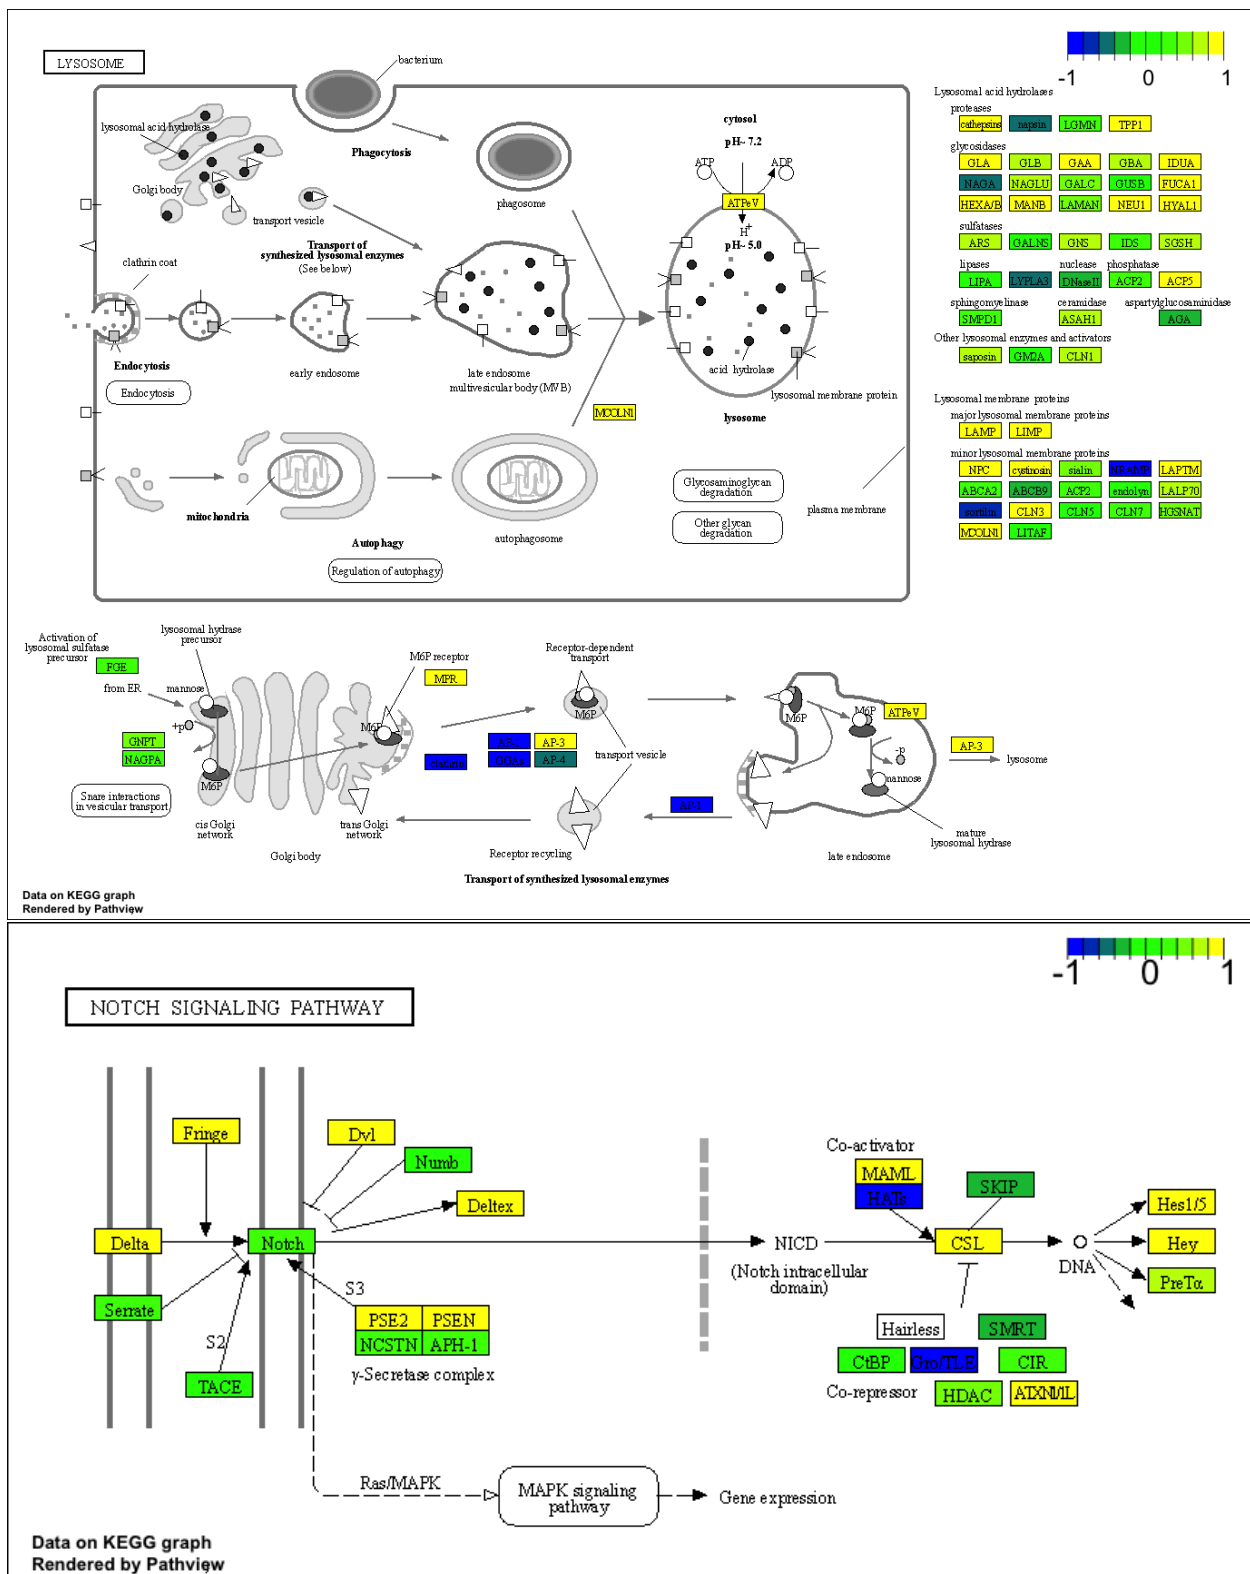
```
## [1] "hsa04640" "hsa04630" "hsa00140" "hsa04142" "hsa04330"
```

Let's pass these 5 IDs to the pathview() function, which will give us a combined output.

```
pathview(gene.data=foldchanges, pathway.id=keggresids, species="hsa", low="blue", mid="green", high="yellow")
```







Q. Can you do the same procedure as above to plot the pathway figures for the top 5 down-regulated pathways?

